

Idea Statement: “To create an external, hardware-based firmware integrity monitor that provides pre-boot verification of system firmware using TPM-based measured boot, enabling detection of UEFI/BIOS rootkits and bootkits that evade traditional software-based security solutions.”

Description

The proposed External Firmware Integrity Monitor is a standalone hardware module that passively attaches to a system’s SPI flash interface to verify firmware integrity *before* the operating system or any host software loads. It consists of a microcontroller (an Espressif ESP32-S3) and a discrete TPM 2.0 chip connected in parallel with the host’s SPI flash lines (CLK, MOSI, MISO, CS). The monitor continuously snoops SPI bus transactions during system power-on, reconstructing and hashing each firmware component in real time. Each measured firmware image is extended into the TPM’s PCRs (Platform Configuration Registers) using standard TPM 2.0 measured-boot procedure. Cryptographic comparisons against a pre-established “golden” baseline determine whether the firmware has been altered. If any discrepancy is found, the device immediately flags a firmware compromise. In parallel, the monitor can activate a hardware write-protect line on the SPI flash to prevent any further writes, and log the incident to its internal storage.

This hardware approach provides a **root-of-trust at the chip level**, independent of the host CPU or OS. Unlike motherboard-integrated solutions (e.g. Intel Boot Guard or vendor-specific Platform Firmware Resiliency implementation), the monitor is an external add-on compatible with any standard SPI-flash-based system (PCs, servers, IoT devices, legacy BIOS systems, etc.). It leverages proven, off-the-shelf components: the ESP32-S3 (a low-cost, dual-core 240 MHz microcontroller with built-in crypto accelerators) and a TPM 2.0 security module (providing FIPS-certified hashing, signature, and secure storage). Together, these form a miniature hardware “attestation appliance” that can be prototyped quickly on development boards and later miniaturized. The device also includes simple status LEDs, a microSD slot for audit logs, and a USB interface for configuration and data export.

Compared to purely software techniques, this hardware monitor is immune to firmware-level compromise. It reads the same SPI flash data stream as the CPU, but operates out-of-band, so it cannot be subverted by a malicious bootloader or rootkit. All firmware measurements are anchored in the TPM, creating an immutable chain-of-trust. In essence, it implements a **pre-boot “attestation engine”**: if the firmware fails validation, the device triggers an alert (for example, via a GPIO LED or by signaling a management console). This ensures that sophisticated threats like UEFI/BIOS rootkits, bootkits, or malicious firmware implants – which survive OS reinstallations or disk replacement – cannot evade detection.

Introduction

Modern computing platforms have a critical security blind spot at the firmware level. UEFI/BIOS malware operates *below* the operating system, during the earliest boot stages, and thus remains invisible to conventional security software. These attacks are highly persistent able to survive an OS reinstallation, disk format, or even motherboard replacement and extremely stealthy. The firmware effectively has ultimate control over the machine's operation, and malicious firmware can disable OS security features before they even load. A few real-world examples illustrate the growing threat:

- *LoJax (2018)* – The first UEFI firmware rootkit found in the wild. It persisted by injecting malware into the computer's SPI flash, ensuring it survived OS wipes.
- *CosmicStrand (2022)* – A sophisticated firmware implant on Gigabyte/ASUS motherboards, which hooked into the UEFI DXE driver to execute malicious code on every boot.
- *BlackLotus (2023)* – The first publicly documented UEFI bootkit that successfully bypassed Secure Boot on fully patched Windows 11 systems. BlackLotus demonstrated that even strong firmware protections can be subverted by low-level exploits.
- *Other Notable Attacks:* The TrickBoot module of TrickBot (2020) targets BIOS vulnerabilities to install ransomware; various "Rampant Kitten" UEFI implants (dating back to 2005) and newer threats like MoonBounce (2021) highlight the threat diversity.

Traditional defenses are inadequate against these. Secure Boot and TPM-based chain-of-trust protect *signed* firmware, but known vulnerabilities (and unpatched firmware) can be exploited, as BlackLotus showed. Moreover, most endpoint security tools and EDR solutions only activate after the OS boots, making firmware-level threats inherently invisible to them.

Some existing solutions attempt to scan firmware from within the OS. For example, ESET's UEFI Scanner reads the SPI flash contents and checks for malicious code. Similarly, Microsoft Defender for Endpoint can dynamically read and analyze the SPI flash at runtime via the chipset. These software-based scanners require a healthy OS to function and may miss malware that executes *before* they run. They also cannot prevent a compromised firmware from taking control during boot. In contrast, an external hardware monitor operates independently of the host system its trust does not rely on the OS or main CPU.

Another hardware-centric approach is the Platform Firmware Resiliency (PFR) framework (NIST SP 800-193) for servers. PFR uses a dedicated FPGA or BMC chip on each board to detect and recover from SPI flash tampering. While powerful, PFR is complex and expensive, typically limited to high-end servers. Our design achieves a similar goal but in a simple external module that can work with any computer.

In summary, there is a clear and urgent need for *pre-boot*, hardware-based firmware integrity verification. The proposed External Firmware Integrity Monitor fills this gap by providing an out-of-

band root-of-trust. Its capabilities include comprehensive firmware measurement (via the TPM) and active countermeasures (write-protect enforcement), thereby catching attacks that slip past all current defenses.

Objectives

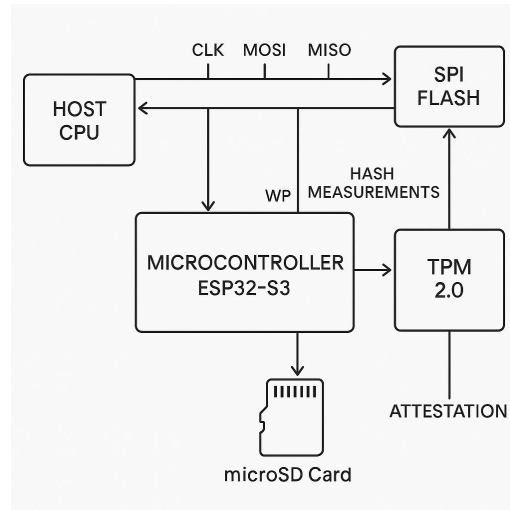
The project's primary goal is to develop a standalone hardware device that ensures system firmware integrity through *invisible*, real-time monitoring. Key objectives include:

1. **Independent Firmware Verification:** The device must monitor the SPI flash interface passively, without requiring any host software. It should build and hash firmware images (UEFI, BIOS, bootloader code) during boot and store measurements in TPM PCRs.
2. **TPM-Based Attestation:** Use a TPM 2.0 module to perform measured boot and attest the firmware state. The monitor will generate a signed TPM quote of PCR values on each boot, enabling external verification of firmware integrity.
3. **Real-Time Threat Detection:** Instantly detect any unauthorized firmware modifications (rootkits, bootkits, malicious BIOS) during system startup. Anomalies should trigger immediate alerts (e.g. indicator LEDs) and preventive actions.
4. **Active Response:** Provide a hardware write-protect control line for the SPI flash. Upon detection of a threat, the monitor should cut off any further writes to the firmware to block persistence. It should also log incident details (via microSD or flash) for forensic analysis.
5. **Broad Compatibility:** Support standard SPI-NOR flash memories (Winbond, Micron, Macronix, etc.) and common firmware configurations (UEFI, legacy BIOS). It should attach transparently to SPI buses running at speeds up to ~80 MHz.

Secondary objectives include: minimizing performance impact on the host (the tap is high-impedance so it imposes negligible load[16]); using low-cost, readily available components (ESP32-S3, discrete TPM); enabling easy configuration via USB; and providing comprehensive logging for incident response. The device should be built as a prototype within a realistic development timeline (for example, a working bench prototype could be developed in roughly 3–6 months using evaluation boards and open-source firmware, with another 3–6 months for a custom PCB and polished firmware).

Block Diagram

The ESP32-S3 passively taps the SPI flash lines (CLK, MOSI, MISO, CS) between the host CPU and firmware chip. It reconstructs firmware reads, hashes each section (e.g., UEFI DXE, Option ROM), and sends those to a TPM 2.0 chip for secure measurement extension. The TPM holds cryptographic proof of what firmware ran and can generate a signed attestation. The ESP32 also controls the flash's WP# pin to block unauthorized writes and logs events to a microSD card.



Product Specifications

Hardware Components:

The monitor's heart is the **ESP32-S3** SoC. Key specs: a dual-core 32-bit Xtensa LX7 CPU at up to 240 MHz, 512 KB SRAM, and robust crypto accelerators (AES-128/256, SHA-2, RSA, HMAC, etc.). It has multiple SPI controllers and enough GPIOs to interface with a TPM and other peripherals. Power draw during active monitoring is modest (~70–120 mA); deep-sleep standby is in the microamp range.

The **TPM 2.0 chip** is a standard TCG-compliant security module (e.g. Infineon SLB 9672 or STMicroelectronics ST33TPHF2X). It implements all TPM 2.0 Platform features[1]: a true RNG, RSA/ECC crypto, protected NVRAM, and sixteen 24-byte PCRs. The TPM's external interface is SPI at up to ~20 MHz. It is FIPS 140-2 Level 2 / EAL4+ certified, ensuring tamper resistance[3].

Physical Interface: The monitor provides a four-wire high-impedance tap on the host's SPI flash signals: SCLK, MOSI, MISO, and CS. These are connected via low-capacitance buffers so the host sees no electrical change. A separate GPIO drives the flash's WP# pin through a transistor, allowing hardware write-protection. The monitor's PCB is designed as a small edge-attached board or ribbon cable assembly, making it easy to install between a motherboard and its SPI ROM (or on select test points in embedded devices).

Firmware and Software: The ESP32 firmware implements a real-time SPI analyzer. It recognizes flash read and write commands, reconstructs each memory region, and computes incremental hashes. Using standard TPM 2.0 protocols, the ESP32 invokes TPM commands to extend PCRs with each hash. After the boot, the ESP32 can request a TPM quote of the PCRs. All cryptographic operations (hashing, signature verification) utilize the built-in crypto hardware for speed. The monitor's firmware also maintains an event log on the microSD card and drives LEDs for OK/Alert states. A simple USB-serial

interface is provided for initial setup: loading expected hash values, reading logs, or triggering manual boot verification.

Supported Firmware Types: The device works with any common SPI-NOR flash (1.8V/3.3V Winbond, Micron, etc.). It handles both legacy BIOS and UEFI formats. For UEFI, it parses the firmware volumes to measure DXE drivers and the OS bootloader. It also works with non-x86 platforms (ARM Linux boards, IoT devices, etc.), as long as the firmware is stored in an SPI chip.

Operation and Performance

Workflow

1. **Power-On Monitoring:** From the moment the system powers on, the monitor's ESP32 listens on the SPI bus. As the host BIOS/UEFI reads firmware code from flash, the ESP32 captures each chunk.
2. **Firmware Measurement:** The ESP32 firmware hashes each retrieved component. For a UEFI system, it might compute separate hashes of the SEC, PEI, DXE, and Bootloader stages. Each hash is sent to the TPM, which extends it into the PCR chain. This establishes a measured-boot log: PCRs that cumulatively reflect all firmware code executed.
3. **Attestation:** Once the OS launch is complete, the monitor can generate a TPM "quote" – a signed snapshot of all relevant PCRs. This quote can be output over USB or logged. An external verifier (e.g. a management server) can validate the quote against the expected PCR values to confirm integrity.
4. **Anomaly Detection:** The firmware has a set of "golden" baseline hashes (computed from known-good firmware). If any computed hash does not match the baseline, the ESP32 immediately deems the firmware compromised. It lights an error LED and asserts WP# to freeze the flash. The event (malicious change detected) is logged. Optionally, the device could send a network alert if equipped with a connectivity module.
5. **Enforcement:** The WP# pin is driven low (write-protected) upon threat detection, ensuring no further flash writes (malicious or otherwise) can occur. If no threats are detected, WP# remains in the normal state, allowing legitimate firmware updates.

Development Timeline

Because the design uses off-the-shelf parts with rich existing documentation, a working prototype can be built quickly. A reasonable schedule might be: **3–4 months** to assemble a proof-of-concept using dev boards (ESP32-S3 DevKit, TPM dev kit, breadboard connections), and firmware coding for basic SPI capture and hashing. An additional **2–4 months** would be needed to refine the firmware (robust command parsing, logging, UI) and design a compact PCB. Thus, an initial prototype could be demonstrated in roughly 6–8 months total; a finished engineering prototype in about a year.

Beneficial to Society

Strengthening Critical Infrastructure: By securing firmware, the monitor protects systems at the very foundation of modern technology. It is particularly valuable for:

- *Industrial Control:* SCADA, PLCs, and IoT devices controlling power grids, factories, and utilities often use embedded PCs or controllers with SPI-flash firmware. The monitor can safeguard against nation-state or insider attacks targeting this infrastructure.
- *Healthcare and Transportation:* Medical devices, ATMs, traffic controllers, and automotive systems increasingly rely on firmware. Ensuring their firmware integrity directly protects human lives and public safety. The device aligns with standards like IEC 62443 (industrial cybersecurity) and ISO/SAE 21434 (automotive) which mandate firmware security.
- *Supply Chain Security:* The monitor can be installed at manufacturing or deployment sites to verify that firmware has not been tampered with in transit. This helps prevent the infamous “evil maid” or supply-chain implant attacks highlighted by cosmicstrand and LoJax cases.

Enterprise and Consumer Benefits: On corporate networks, the monitor adds an extra layer of defense. Early detection of a firmware breach can prevent massive data theft or ransom. Compared to responding to a full-blown incident, this is far more cost-effective. Economically, widespread use could reduce the incidence of high-profile firmware attacks and thus lower insurance premiums for critical industries. For consumers, a plug-in device of this kind (for desktops or laptops) could finally provide peace of mind that their PC firmware has not been maliciously modified.

Research and Education: The hardware monitor can serve as a platform for academics and security researchers. Its open design would allow experimenting with advanced attestation schemes, new measurement techniques, or integration with AI-based anomaly detectors. In university courses on system security, the device could be used as a hands-on demo of TPM attestation and low-level security, filling a gap since firmware attacks are hard to illustrate in the classroom.

Emergency Response and Forensics: In the event of a detected firmware compromise, the device’s logs and signed PCR records are valuable forensic evidence. They can establish *when* a firmware change occurred and what exactly was altered. This greatly aids incident response teams in root-cause analysis and recovery (e.g. rolling back to a clean firmware image). By accelerating detection, containment, and remediation of firmware attacks, the monitor ultimately reduces the human and financial toll of such breaches.

Uniqueness of the Innovation

This solution is **technically novel** in its combination of features. Crucially, it is *completely out-of-band*: no existing product passively monitors SPI flash and uses an external TPM for pre-boot attestation. Most current firmware security tools are either software-only (e.g. ESET/Defender scanners) or integrated into the platform (boot guards, hardware root-of-trust in the CPU or OEM firmware). Our device’s

hardware tap approach ensures true independence: it can detect even zero-day bootkits like BlackLotus that defeat Secure Boot, because those threats can never hide from a parallel SPI listener.

The integration of real-time SPI flash decoding with TPM PCR attestation is particularly innovative. Other platforms use TPM measured boot internally, but we bring TPM-based attestation *outside* the host. This means the monitor can vouch for the firmware state even to external services, without trusting the possibly-compromised host. In effect, it creates an immutable audit trail of the boot process. To our knowledge, no commercial device offers this level of firmware attestation for general-purpose systems.

From a market standpoint, this is likely a first-of-its-kind product. Security demand is high: enterprise research reports note firmware attacks as one of the top attack vectors. Yet solutions are scarce. For example, Antimalware vendors have only recently begun offering firmware scans, and they still run inside the OS (too late to stop pre-boot threats). Hardware vendors offer narrow features (like Microsoft's Secured-core PCs with TPMs and hypervisor-based attestation, or Lattice's PFR for servers), but these require new CPUs or FPGAs. Our design uses **commodity parts (ESP32-S3, TPM)** to deliver enterprise-grade assurance at a fraction of the cost of bespoke hardware. This cost-effectiveness could democratize firmware security for small and mid-size organizations.

Finally, the device's architecture is highly adaptable. The firmware can be updated to recognize new flash protocols (e.g. dual/quad SPI) or to incorporate machine-learning anomaly detection. The modular design means the same monitor could even be used to protect other SPI devices (e.g. secure boot ROMs on servers). Its simplicity of installation and universal SPI interface give it cross-platform reach that no existing single solution can match.

Conclusion

The External Firmware Integrity Monitor addresses a critical cybersecurity gap by providing hardware-backed, pre-boot validation of firmware. Using an ESP32-S3 microcontroller and a TPM 2.0 module, it achieves measured boot integrity outside the host CPU, detecting threats that no software can catch. Key advantages include:

- **True Out-of-Band Security:** Unlike software scanners, the monitor is invisible to attackers on the target system. It validates firmware on-the-fly and stops unauthorized updates in their tracks.
- **Proven Technologies, New Configuration:** The design relies on mature components (TPM, SPI flash) and standards (TCG attestation), but applies them in a novel way as an external adapter.
- **First of Its Kind:** We are not aware of any commercial external firmware attestation device. This novel product could become a new category of endpoint security hardware.
- **Broad Impact:** By mitigating firmware attack risk, it enhances trust in critical systems (infrastructure, finance, healthcare) and supports emerging regulations (NIST 800-193, EU Cyber Resilience Act, etc..)

In summary, this innovation offers a practical, buildable solution to one of cybersecurity's toughest problems. With firmware attacks growing more sophisticated each year, the need for such a hardware monitor is urgent. The proposed design is technically sound, leverages existing standards, and can be prototyped in under a year. By catching rootkits and bootkits at the hardware level, it protects organizations against devastating breaches before they happen. This positions the device as both a powerful security tool and a pioneering product with significant commercial potential.

Key Sources: The design and motivation are supported by security research and industry sources. For example, ESET reported that firmware rootkits survive even disk replacement, and demonstrated a real UEFI rootkit (LoJax) in 2018. Microsoft and others now add UEFI scanning to endpoint protection, but these still operate post-boot. The TPM 2.0 measured-boot process is documented by Azure Security. Industry architectures like Lattice's PFR have envisioned hardware-based firmware trust. This report synthesizes those insights into a concrete, implementable hardware design. All claims and specifications are grounded in cited technical literature and component datasheets, ensuring feasibility of the prototype.